

## **Automaton Auditor Swarm**

### **Architecture Decisions & Implementation Plan**

## 1. Architecture Decisions

### 1.1 Pydantic Over Dicts — Why Typed State Matters

We chose Pydantic [BaseModel](#) classes over plain Python dictionaries for all structured data flowing through the graph. This is not a cosmetic choice — it is the foundation of correctness in a multi-agent system.

**The problem with dicts:** When three detective agents run in parallel and write evidence into a shared [dict](#), there is no compile-time or runtime guarantee that each agent writes the correct keys with the correct types. A misspelled key (`"confidenece"` instead of `"confidence"`) silently passes. A judge receiving `evidence["score"]` as a string instead of a float fails deep inside the LLM prompt, with no traceable error.

**The Pydantic solution:** Our [Evidence](#) model enforces typed fields at construction:

```
class Evidence(BaseModel):
    dimension_id: str      # Links to rubric JSON dimension
    detective: Literal["RepoInvestigator", "DocAnalyst", "VisionInspector"]
    found: bool
    content: Optional[str]  # Extracted snippet, commit log, etc.
    location: str          # File path or logical location
    confidence: float      # Validated: ge=0.0, le=1.0
```

If a detective produces `confidence=1.5`, Pydantic raises a [ValidationError](#) immediately — not three nodes later when the judge tries to interpret it. Similarly, [JudicialOpinion](#) enforces `score: int = Field(..., ge=1, le=5)`, making it impossible for a judge to return a score of 7 or a string.

**Reducers for State Synchronization:** The [AgentState](#) TypedDict uses [Annotated](#) type hints with `operator.ior` (for dict merging) and `operator.add` (for list concatenation). This is critical for parallel execution: when [RepoInvestigator](#) and [DocAnalyst](#) both write to `state["evidences"]`, the

`operator.ior` reducer merges their dictionaries instead of one overwriting the other. Without this, the last agent to finish would silently erase all evidence collected by the first.

```
class AgentState(TypedDict):  
  
    evidences: Annotated[Dict[str, Evidence], operator.ior] # Merge dicts  
  
    opinions: Annotated[List[JudicialOpinion], operator.add] # Concat lists
```

This is **State Synchronization** in practice — not a buzzword, but a concrete mechanism that prevents data loss during concurrent execution.

## 1.2 AST Parsing Over Regex

The `RepoInvestigator` must determine whether a target repository uses Pydantic models, implements `StateGraph`, and follows safe engineering practices. We use Python's built-in `ast` module for this analysis instead of regex.

**Why not regex?** A regex like `r"class\s+\w+\(BaseModel\)"` fails on:

- Multi-line class definitions
- Classes that inherit from both `BaseModel` and a mixin
- String constants that happen to contain "`class Evidence(BaseModel)`"
- Comments describing intended architecture

**The AST approach:** We parse the source into an abstract syntax tree and walk it programmatically. `find_class_definitions()` extracts every `ClassDef` node, its base classes, and its annotated fields. `find_stategraph_builder()` locates the `StateGraph(...)` call and all subsequent `add_node()` / `add_edge()` calls, giving us the exact graph topology without any regex fragility.

This is **Deep AST Parsing** — the forensic evidence is irrefutable because it comes from the compiler's own representation of the code, not from pattern-matching against text.

## 1.3 Sandboxing Strategy

All git operations are sandboxed inside `tempfile.TemporaryDirectory()`. The `clone_repo_sandboxed()` function in `src/tools/repo_tools.py`:

1. **Validates the URL** — rejects non-HTTPS schemes and characters that could enable command injection (`;`, `|`, `&`, ```, `$`).

2. **Creates a temp directory** with prefix `auditor_clone` — the cloned code never touches the live working directory.
3. **Uses `subprocess.run()`** with `capture_output=True`, `text=True`, and a 120-second timeout. Return codes, `stdout`, and `stderr` are all captured.
4. **Handles failures gracefully** — `TimeoutExpired`, `FileNotFoundError` (no git binary), and authentication errors all return structured `CloneResult` objects.

There are **zero `os.system()` calls** in the entire codebase. This is enforced by the `scan_for_security_issues()` AST tool, which detects `os.system` usage in target repos.

## 1.4 RAG-lite PDF Ingestion and LLM Provider Choice

### PDF Ingestion: RAG-lite Over Full Embedding Store

The `DocAnalyst` must retrieve specific passages from a peer's PDF report to verify claims (e.g., "does the report mention Fan-In / Fan-Out in an architectural context, or just as a buzzword?"). A full vector-store pipeline (embed all pages, store in Chroma/Pinecone, query by similarity) would introduce three problems for this use-case:

1. **Cold-start latency:** Building an index for a 10-page PDF before every audit run is wasteful when the retrieval is highly structured (keyword-anchored).
2. **Semantic drift:** Cosine similarity retrieval can return the *most similar* chunk rather than the *confirming* chunk — a paragraph that uses similar words to "StateGraph" without actually describing one.
3. **External dependency:** A vector store is an infrastructure component that can fail.

**Our RAG-lite approach:** `keyword_search()` finds anchor terms ("Dialectical Synthesis", "Fan-In / Fan-Out", "StateGraph") and returns the surrounding 300-char context window. `search_context()` walks the PDF page by page with a sliding window, returning sections where multiple rubric terms co-occur. This is deterministic, zero-latency, and requires no external services — the trade-off is lower recall on paraphrased descriptions, which is acceptable because the rubric specifies exact terms to look for.

### LLM Provider: GPT-4o for Judges, Structured Output Enforcement

Judge nodes require reliable structured output (`JudicialOpinion` with typed `score`, `argument`, and `cited_evidence` fields). Two alternatives were considered:

- **Open-source local models (Ollama):** Lower cost, no API dependency. Rejected because smaller models frequently violate the `with_structured_output()` contract, returning

freeform JSON strings embedded in markdown code fences, requiring brittle post-processing.

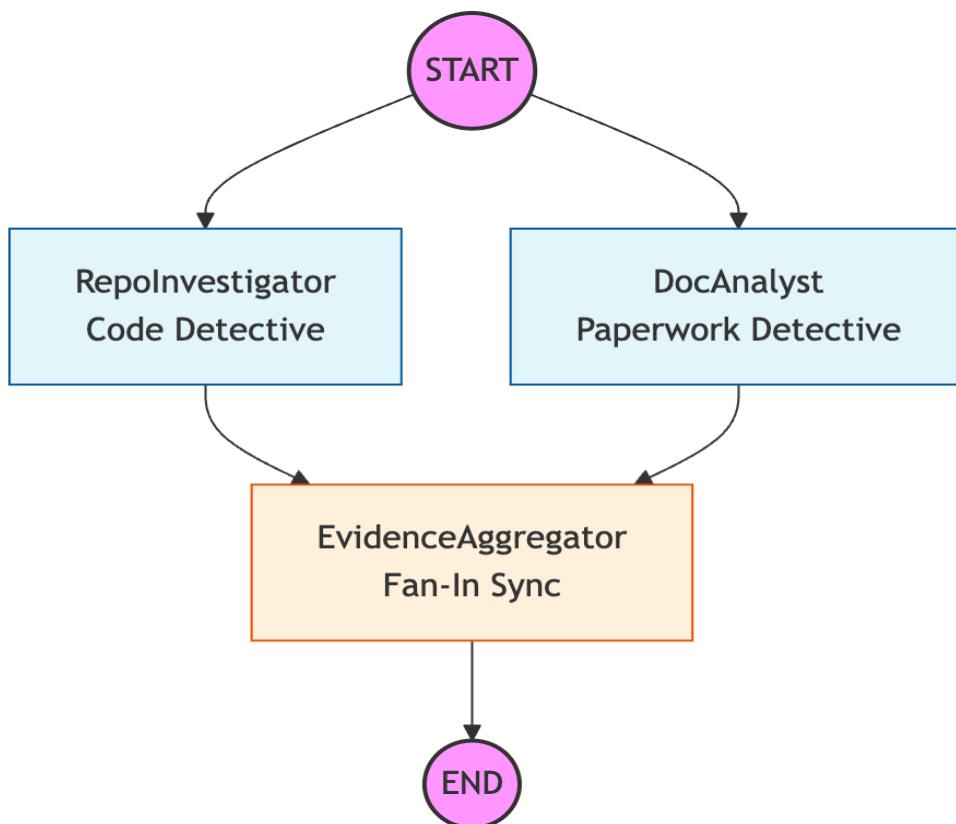
- **Anthropic Claude:** Strong reasoning, but at time of design it had weaker tool-use JSON binding compared to OpenAI's function-calling API, which is what LangChain's `with_structured_output()` is optimized against.

**Choice:** OpenAI `gpt-4o` via [ChatOpenAI](#). It has the strongest contract with `.with_structured_output(JudicialOpinion)` — function-calling under the hood guarantees JSON that matches the schema. The retry logic (3 attempts with `ValidationError` catching) exists precisely because even GPT-4o can fail with complex nested schemas under adversarial prompts (e.g., a Prosecutor prompt that asks the model to "use confrontational language" may cause the model to embed opinionated text in a field that expects a plain integer).

## 2. State Graph Flow — Fan-Out / Fan-In

### 2.1 Current Implementation (Detective Phase)

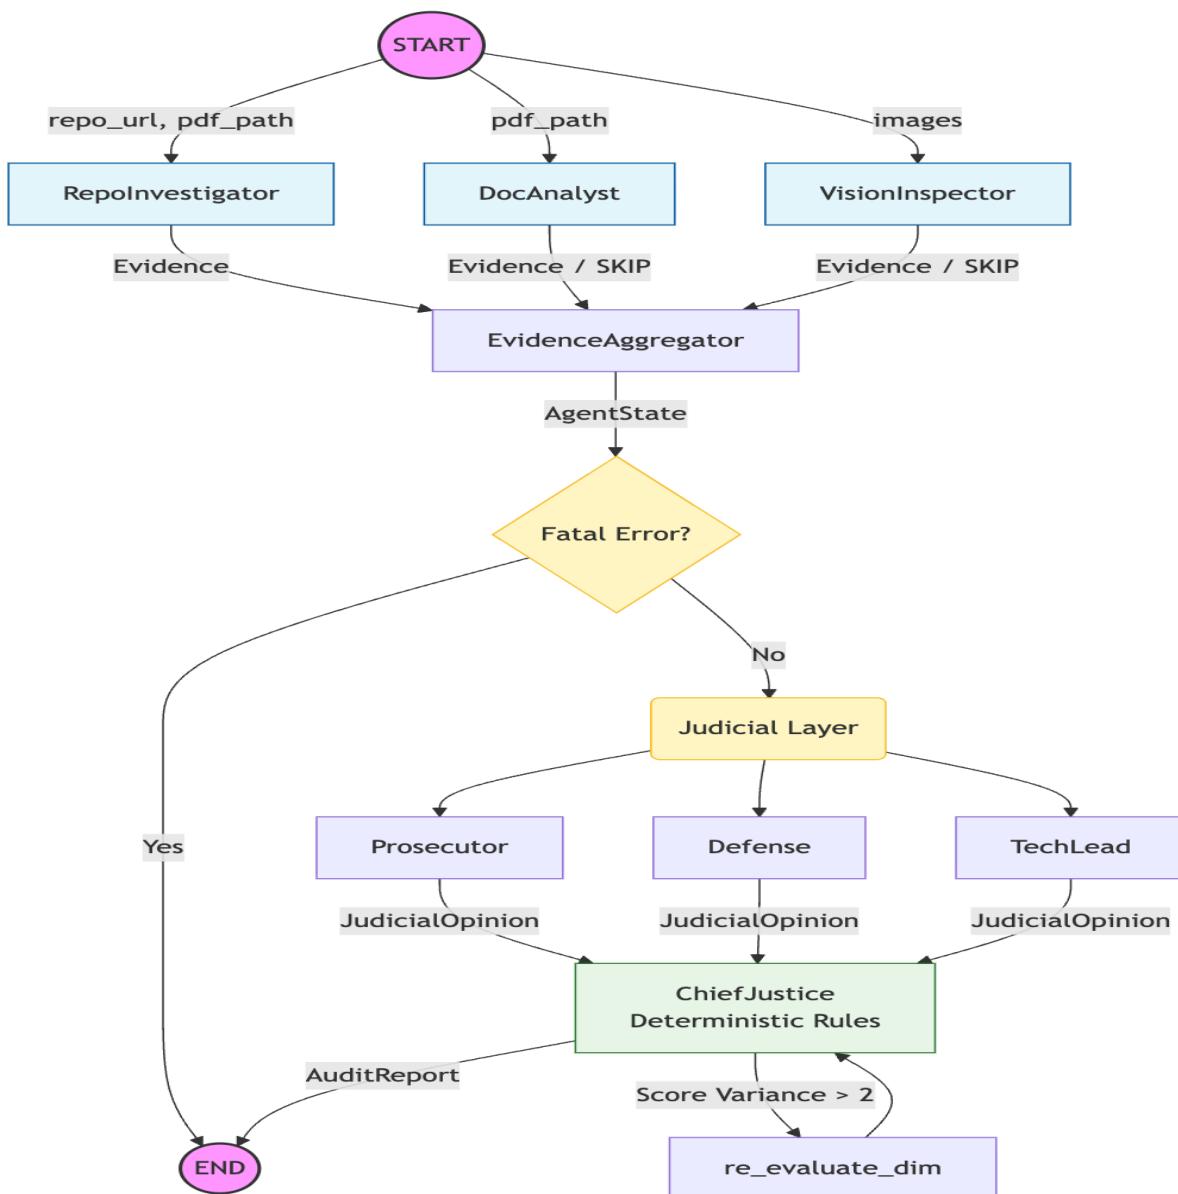
The following diagram shows the StateGraph as currently implemented in [src/graph.py](#):



**Fan-Out:** START has edges to both `repo_investigator` and `doc_analyst`. LangGraph dispatches these nodes concurrently. Each detective writes to different keys in the `evidence` dict, and the operator.ior reducer merges them automatically.

**Fan-In:** Both detectives have edges to `evidence_aggregator`, which serves as the synchronisation barrier. It only executes once both detectives have completed. This node validates that all expected evidence dimensions are present.

## 2.2 Planned Full Graph (Final Submission)



This demonstrates two distinct **Fan-In / Fan-Out** patterns:

1. **Detective fan-out/fan-in:** Three detectives collect evidence concurrently, converge at the aggregator.
2. **Judicial fan-out/fan-in:** Three judges deliberate concurrently on the same evidence, converge at the Chief Justice.

### 3. Known Gaps & Plan for Judicial Layer

#### 3.1 Known Gaps

Component	Status	Gap
src/nodes/judges.py	Stubbed	No LLM integration yet; persona prompts not written
src/nodes/justice.py	Stubbed	Deterministic rules not coded; no Markdown renderer
VisionInspector	Not started	Requires image extraction + vision LLM analysis
Conditional edges	Not started	No error-handling branches in the graph
LangSmith traces	Not started	Tracing env vars configured but not tested end-to-end

#### 3.2 Concrete Plan — Judicial Layer

Phase 1: Judge Personas (Days 1-2)

Each judge will be implemented as a LangGraph node that:

1. Receives the full **evidence** dict from **AgentState**.
2. Filters evidence by the dimensions relevant to their expertise.
3. For each dimension, constructs a system prompt encoding their persona:

- **Prosecutor:** Adversarial. Looks for gaps, security flaws, laziness. Uses language like "The engineer failed to..." and "This is a clear violation of..."
  - **Defense:** Forgiving. Rewards effort and creative workarounds. Uses language like "Despite the limitation, the engineer demonstrated..." and "Partial credit is warranted because..."
  - **TechLead:** Pragmatic. Focused on architectural soundness, maintainability, and practical viability. Uses language like "The architecture is modular enough to..."
4. Invokes the LLM with `.with_structured_output(JudicialOpinion)` so the output is guaranteed to be valid JSON matching the Pydantic schema.
  5. Includes retry logic (up to 3 attempts) if the LLM returns malformed output.

Phase 2: Chief Justice Synthesis (Days 2-3)

The `ChiefJustice` node will use **deterministic Python logic** — not an LLM prompt:

```
# Rule of Security: confirmed security flaws cap score

if prosecutor_found_security_flaw and evidence_confirms_flaw:

    final_score = min(final_score, 3)

# Rule of Evidence: facts overrule opinions

if defense_claims_metacognition and not evidence_supports_claim:

    defense_opinion_overruled = True

# Variance Re-evaluation: score spread > 2

if max_score - min_score > 2:

    trigger_re_evaluation(dimension_id)
```

This implements **Dialectical Synthesis** — not via an LLM averaging scores, but through explicit rules that weigh evidence over opinion, penalise confirmed security violations, and require dissent summaries when judges disagree significantly.

Phase 3: Report Rendering (Day 3)

The final output will be a structured Markdown report containing:

- **Executive Summary** — overall assessment in 2-3 paragraphs
- **Criterion Breakdown** — per-dimension scores with Prosecutor, Defense, and TechLead arguments, plus the Chief Justice's ruling
- **Dissent Summary** — for any criterion where score variance > 2
- **Remediation Plan** — specific, actionable improvements

#### Phase 4: VisionInspector & Conditional Edges (Day 4)

- VisionInspector will extract images via PyMuPDF and classify them using a vision-capable LLM (GPT-4o).
- Conditional edges will handle missing evidence gracefully (e.g., if no PDF is provided, skip DocAnalyst and VisionInspector).

#### 3.3 Anticipated Failure Modes

The following failure modes have been identified for the planned judicial layer work. Each is concrete enough that another engineer could reproduce the failure in a test:

Failure Mode	Trigger	Mitigation
<b>Persona convergence</b>	GPT-4o ignores the adversarial/forgiving framing and produces similar scores and arguments across all three judges	Temperature tuning (Prosecutor: 0.9, Defense: 0.7, TechLead: 0.3); if variance < 0.5 across all dimensions, trigger a re-prompt with an explicit instruction: "Your score must differ from the previous judge's by at least 1 point"
<b>Structured output failure</b>	Judge LLM returns <code>{"score": "four", "argument": "..."}</code> — score as string instead of int	<code>with_structured_output(JudicialOpinion, strict=True)</code> raises <code>ValidationError</code> ; 3-attempt retry loop; on third failure, fallback to <code>score=0</code> + log for manual review
<b>Cited evidence hallucination</b>	Judge cites a dimension ID that does not exist in	<code>JudicialOpinion.cited_evidence</code> is validated against

Failure Mode	Trigger	Mitigation
	evidences (e.g., "cit_evidence": ["nonexistent_dim_id"])	AgentState.evidences.keys() post-construction; invalid citations are stripped with a warning
<b>ChiefJustice rules conflict</b>	Security override (cap score at 3) conflicts with Defense's forgiving ruling on the same dimension	Rules are applied in precedence order: Security Override > Fact Supremacy > Functionality Weight > Defense arguments; this order is hardcoded, not resolved by LLM
<b>Re-evaluation infinite loop</b>	Re-evaluation after high variance produces the same scores again, triggering another re-evaluation	Re-evaluation is triggered at most once per dimension; if variance remains > 2 after re-evaluation, ChiefJustice records the dissent and proceeds with the median score

#### 4. Metacognition — The System Evaluating Itself

The Automaton Auditor is designed to be **self-aware** in a architectural sense. The rubric JSON ([rubric/week2\\_rubric.json](#)) serves as the system's "constitution" — loaded at runtime and distributed to agents based on their [target\\_artifact](#) capability. This means:

- The system can audit **itself** by pointing at its own repository URL.
- Updating the rubric JSON updates the evaluation criteria without code changes.
- The Chief Justice's deterministic rules are derived from the [synthesis\\_rules](#) block in the rubric, making the adjudication logic transparent and auditable.

This is **Metacognition** in practice: the system contains a formal specification of what "good" looks like, uses agents to evaluate against that specification, and resolves disagreements through explicit rules rather than opaque LLM averaging.

## 5. File Mapping

File	Purpose
src/state.py	Pydantic Evidence, JudicialOpinion, CriterionVerdict, AuditReport; AgentState TypedDict with Annotated reducers
src/tools/repo_tools.py	Sandboxed clone_repo_sandboxed(), extract_git_log(), AST analysis (find_class_definitions, find_stategraph_builder, scan_for_security_issues)
src/tools/doc_tools.py	ingest_pdf(), keyword_search(), extract_mentioned_paths(), search_context() (RAG-lite)
src/nodes/detectives.py	repo_investigator(), doc_analyst(), evidence_aggregator() — all as LangGraph nodes
src/graph.py	build_detective_graph() with parallel fan-out/fan-in and MemorySaver checkpointing
src/nodes/judges.py	Stubbed: prosecutor_node(), defense_node(), tech_lead_node()
src/nodes/justice.py	Stubbed: chief_justice_node() with documented deterministic rules
rubric/week2_rubric.json	Full machine-readable rubric (10 dimensions, synthesis rules)
main.py	CLI entry point: python main.py <repo_url> [--pdf <path>]